

Variational minimisation

This notebook analytically and numerically explores variational minimisation, demonstrating quantum gradient descent and quantum natural gradient.

Contents:

- *Analytic*
- *Gradient descent*
- *Natural gradient*


Tyson Jones, 2024
Department of Materials, University of Oxford
tyson.jones.input@gmail.com

```
Import["https://qtechtheory.org/questlink.m"];  
CreateDownloadedQuESTEnv[];
```

Analytic

Variational quantum algorithms make use of a parameterised “ansatz” circuit $U(\vec{\theta})$ acting upon a fixed input state $|in\rangle$, in order to produce parameterised states $|\psi(\vec{\theta})\rangle$. For example, consider this three-qubit ansatz with parameters $\vec{\theta} = \{a, b\}$.

```
nQb = 3;  
u = Circuit[ H0 S0 H1 S1 H2 S2 C0[X2] Rz2[a] Rz1[b] C2[X1] ];  
DrawCircuit[u]
```



Let's use a fixed input state $|in\rangle = |0\rangle$.

```
in = UnitVector[2nQb, 1]  
{1, 0, 0, 0, 0, 0, 0, 0}
```

We can express our ansatz state $|\psi(\vec{\theta})\rangle$ analytically as a function of a and b

```
 $\psi = \text{CalcCircuitMatrix}[u] . \text{in} // \text{Simplify}$ 
```

$$\left\{ \frac{e^{-\frac{1}{2}i(a+b)}}{2\sqrt{2}}, -\frac{e^{-\frac{1}{2}i(a+b)}}{2\sqrt{2}}, \frac{ie^{-\frac{1}{2}i(a-b)}}{2\sqrt{2}}, \right. \\ \left. -\frac{ie^{-\frac{1}{2}i(a-b)}}{2\sqrt{2}}, -\frac{e^{\frac{1}{2}i(a+b)}}{2\sqrt{2}}, -\frac{e^{\frac{1}{2}i(a+b)}}{2\sqrt{2}}, \frac{ie^{\frac{1}{2}i(a-b)}}{2\sqrt{2}}, \frac{ie^{\frac{1}{2}i(a-b)}}{2\sqrt{2}} \right\}$$


An experimentalist could freely change the values of parameters a and b , smoothly changing the output quantum state. For instance, here's how the fidelity with the $|+\rangle$ state would change as b is varied (incidentally, it is independent of a).

```
plus = ConstantArray[ $\frac{1}{\sqrt{2^{n_{\text{Qb}}}}}$ ,  $2^{n_{\text{Qb}}}$ ];
```

```
fid = Simplify[Abs[ $\psi \cdot \text{plus}$ ]2, {a, b} ∈ Reals]
```

$$\frac{1}{16} \text{Abs}[-i + e^{ib}]^2$$

```
Plot[fid, {b, -2 π, 2 π}, AxesLabel → {"b", " $|\langle \psi(a,b) | + \rangle|^2$ "}]
```



We are often interested in the observables of some operator, like a Hamiltonian, natively expressed as a Pauli string. Here's one I just made up:

$$h = X_0 Y_1 + 2 Z_1 Z_2 - 3 Y_0 Y_2 ;$$

and here is now it looks as a 3-qubit Z-basis matrix:


```
hM = Normal @ CalcPauliExpressionMatrix[h];
MatrixForm[hM]
```

$$\begin{pmatrix} 2 & 0 & 0 & -i & 0 & 3 & 0 & 0 \\ 0 & 2 & -i & 0 & -3 & 0 & 0 & 0 \\ 0 & i & -2 & 0 & 0 & 0 & 0 & 3 \\ i & 0 & 0 & -2 & 0 & 0 & -3 & 0 \\ 0 & -3 & 0 & 0 & -2 & 0 & 0 & -i \\ 3 & 0 & 0 & 0 & 0 & -2 & -i & 0 \\ 0 & 0 & 0 & -3 & 0 & i & 2 & 0 \\ 0 & 0 & 3 & 0 & i & 0 & 0 & 2 \end{pmatrix}$$

By studying the expectation value of this observable upon states output from our ansatz circuit, we are effectively probing a parameterised manifold of the observable space.

```
v = Conjugate[ψ] . hM . ψ;
v = FullSimplify[v, {a, b} ∈ Reals]
-Cos[b] + 3 Sin[a] Sin[b]
```

```
Plot3D[v, {a, -2 π, 2 π}, {b, -2 π, 2 π}, AxesLabel → {"a", "b", "⟨ψ|h|ψ⟩"}]
```



Often we are interested in the minimum eigenvalue of the observable. If our operator \mathbf{h} is a Hamiltonian, this is the ground-state energy.

```
Min @ Eigenvalues @ hM
-√14
```


Alas, our parameterised circuit generates only a strict subspace of states $|\psi(\vec{\theta})\rangle$, which is unlikely to contain the true ground-state. Indeed, the lowest our circuit above can produce is:

```
MinValue[v, {a, b}]
-√10
```

In principle, we can add more parameters to our circuit and increase the size of our accessible subspace. Here, we'll add introduce additional controlled rotations with parameters **c** and **d**...

```
u = Join[u, { C0[Rx1[c]], C1[Rz2[d]] }];
```

```
DrawCircuit[u]
```



```
ψ = CalcCircuitMatrix[u] . in;
```

```
v = Conjugate[ψ] . hM . ψ;
```

```
v = FullSimplify[v, {a, b, c, d} ∈ Reals]
```

$$\frac{1}{2} \cos\left[\frac{c}{2}\right] \left(-3 \cos[a+b] - 2 \cos\left[b - \frac{d}{2}\right] + 3 \cos[a-b+d] + 4 \cos[b] \sin\left[\frac{c}{2}\right] \right)$$

which enables producing states a little closer to the groundstate at $-\sqrt{14} \approx -3.74$

```
NMinimize[v, {a, b, c, d}] // Chop
```

```
{-3.63226, {a → -0.321751, b → 0, c → -0.855463, d → -2.49809}}
```

Beware that a larger parameterised observable space is slower to explore! With only a few more parameters and qubits, our variational system also would become too large to study analytically like we do here, and we would have to resort to numerical study, as we do below.

Gradient descent

Let's consider a random **7-qubit 30-term** Hamiltonian

```
nQb = 7;
```

```
h = GetRandomPauliString[nQb, 30, {-1, 1}]
```


```
0.8042 X1 X3 + 0.588717 X0 X5 X6 Y1 - 0.913716 X2 X3 Y0 Y1 Y4 Y6 +
0.954613 X0 Y2 Y4 Y5 Z1 - 0.0842519 X2 X5 Y0 Y3 Y4 Y6 Z1 + 0.740256 X1 X3 X4 Y5 Z2 -
0.714507 X3 X4 Y1 Y5 Z2 + 0.403638 X6 Y1 Y5 Z2 + 0.675392 X5 Y0 Y4 Z1 Z2 +
0.485387 X2 X5 Y0 Y1 Y6 Z3 + 0.769312 X1 X6 Y0 Y2 Y3 Z4 - 0.027828 X2 X3 X5 Y1 Z0 Z4 +
0.652981 X0 X1 X2 X5 Y6 Z3 Z4 - 0.421427 X2 X5 X6 Y1 Z0 Z3 Z4 + 0.379304 X0 X4 Y1 Y6 Z5 +
0.113742 X1 X2 Y4 Z0 Z5 - 0.972824 X3 Y4 Y6 Z0 Z5 + 0.0600258 X6 Y0 Y1 Z3 Z5 -
0.384336 X1 Z0 Z3 Z5 - 0.366232 X0 X2 X4 Y6 Z1 Z3 Z5 - 0.582251 X0 X1 X2 X6 Y3 Z4 Z5 -
0.833795 X3 X5 Y0 Y1 Y2 Z6 - 0.705635 X1 X2 Y4 Z0 Z3 Z6 - 0.515419 X3 X5 Y1 Z0 Z4 Z6 -
0.867807 X0 X3 Y5 Z1 Z4 Z6 - 0.847075 X0 Y1 Y2 Z3 Z4 Z6 + 0.25922 X0 X1 Y3 Z5 Z6 -
0.928031 Y0 Y3 Y4 Z5 Z6 - 0.250159 Y3 Z1 Z2 Z5 Z6 - 0.269194 X1 X2 X3 Z0 Z4 Z5 Z6
```

```
vMin = CalcPauliStringMinEigVal[h]
```

```
-6.45173
```

Imagine a quantum experimentalist seeks the ground-state energy of this Hamiltonian, and can freely vary **56** parameterised gates in their ansatz circuit $u(\vec{\theta})$ which is applied to initial state $|+\rangle$.

```
u = GetKnownCircuit["HardwareEfficientAnsatz", 3,  $\theta$ , nQb];
DrawCircuit[u]
```



```
n $\theta$  = Max @ Cases[u,  $\theta[i_] \mapsto i$ ,  $\infty$ ]
56
```

We will numerically simulate the experimental process, so we prepare some quantum registers in our simulator.

```
{ $\psi$ ,  $\phi$ , in} = CreateQuregs[nQb, 3];
InitPlusState[in];
```

The experimentalist cannot study an analytic expression of their ansatz state $|\psi(\vec{\theta})\rangle = u(\vec{\theta})|+\rangle$, nor its expectation value, which are exponentially expensive! Instead, they can only measure the observable $\langle E(\vec{\theta}) \rangle = \langle \psi(\vec{\theta}) | h | \psi(\vec{\theta}) \rangle$ at specific values of $\vec{\theta}$.

```
CloneQureg[ $\psi$ , in];
```

```
u /.  $\theta[_] \mapsto$  RandomReal[]
```


```
{Ry0[0.379964], Rz0[0.416136], Ry1[0.516909], Rz1[0.745819], Ry2[0.253023],
Rz2[0.782511], Ry3[0.483642], Rz3[0.358359], Ry4[0.987236], Rz4[0.0847689],
Ry5[0.233927], Rz5[0.606911], Ry6[0.804772], Rz6[0.900195], C0[Z1], C2[Z3],
C4[Z5], C6[Z0], C1[Z2], C3[Z4], C5[Z6], Ry0[0.178302], Rz0[0.764123], Ry1[0.633428],
Rz1[0.176239], Ry2[0.825577], Rz2[0.0903814], Ry3[0.0710499], Rz3[0.325454],
Ry4[0.114324], Rz4[0.71113], Ry5[0.408383], Rz5[0.872261], Ry6[0.0330232],
Rz6[0.172492], C0[Z1], C2[Z3], C4[Z5], C6[Z0], C1[Z2], C3[Z4], C5[Z6], Ry0[0.519215],
Rz0[0.899798], Ry1[0.782472], Rz1[0.0897789], Ry2[0.0161363], Rz2[0.757917],
Ry3[0.979974], Rz3[0.195149], Ry4[0.276458], Rz4[0.448323], Ry5[0.955637],
Rz5[0.427049], Ry6[0.239681], Rz6[0.150582], C0[Z1], C2[Z3], C4[Z5], C6[Z0], C1[Z2],
C3[Z4], C5[Z6], Ry0[0.267701], Rz0[0.17584], Ry1[0.0386318], Rz1[0.196578],
Ry2[0.72915], Rz2[0.375773], Ry3[0.275606], Rz3[0.513447], Ry4[0.6292],
Rz4[0.747971], Ry5[0.0212156], Rz5[0.836306], Ry6[0.0610466], Rz6[0.549128]}
```

```
ApplyCircuit[ψ, %];
CalcExpecPauliString[ψ, h, ϕ]
-0.0822767
```

The parameter space is already too big to randomly sample like this. We are *very* unlikely to randomly choose parameters near the ground-state, especially given the *barren plateaus*.

```
samples = Table[
  CloneQureg[ψ, in];
  ApplyCircuit[ψ, u /. θ[_] :-> RandomReal[]];
  CalcExpecPauliString[ψ, h, ϕ],
  20
]
{0.290054, 0.0354827, 0.54823, 0.242395, 0.0316951, 0.403997,
 0.241329, 0.573669, 0.261465, 0.371843, 0.241047, 0.100007, 0.124764,
 0.381409, -0.168843, 0.565205, 0.418175, -0.1838, 0.287822, 0.337146}
```

```
ListPlot[samples, PlotRange -> {1.1 vMin, .5},
  GridLines -> {{}, {{vMin, Directive[Red, Dashed]}}},
  AxesLabel -> {"sample", "energy"}]
```



The experimentalist could instead employ quantum gradient descent, treating the observable's expectation value as the cost function to be minimised. This requires they can additionally obtain the *gradient* of the observable, $\nabla_{\vec{\theta}} \langle \psi(\vec{\theta}) | h | \psi(\vec{\theta}) \rangle$, at a given position in parameter space. One method to do so is via the parameter shift rule.

Consider a unitary gate of form $U(\theta) = \exp(-i a \theta G)$ with Hermitian generator G , and observable $\langle E(\theta) \rangle = \langle \psi | U^\dagger(\theta) H U(\theta) | \psi \rangle$. The parameter shift rule states that

$\frac{d}{d\theta} \langle E(\theta) \rangle = r [\langle E(\theta + \frac{\pi}{4r}) \rangle - \langle E(\theta - \frac{\pi}{4r}) \rangle]$ where $r = \frac{a}{2} (e_1 - e_0)$, and e_0 and e_1 are the eigenvalues of G . Assuming all ansatz gates depend on unique parameters, this expression gives the full circuit's derivative.

Because our ansatz is composed of $R_y = \exp(-i \theta/2 Y)$ and $R_z = \exp(-i \theta/2 Z)$, which have Y and Z generators with ± 1 eigenvalues, we know that $a = \frac{1}{2}$ and $r = a$, so that

$\frac{d}{d\theta} \langle E(\theta) \rangle = \frac{1}{2} [\langle E(\theta + \frac{\pi}{2}) \rangle - \langle E(\theta - \frac{\pi}{2}) \rangle]$. Our experimentalist only needs to obtain *two* expectation values in order to find the derivative of the expectation with respect to a parameter. They don't need any new circuits - they just sample the ansatz circuit!

```
calcExpecDeriv[h_, in_, u_, vθ_, dθ_, ψ_, φ_] := Module[
  {v1, v2, e1, e2},

  v1 = vθ /. (dθ → v_) → (dθ → v + π / 2);
  v2 = vθ /. (dθ → v_) → (dθ → v - π / 2);

  CloneQureg[ψ, in];
  ApplyCircuit[ψ, u /. v1];
  e1 = CalcExpecPauliString[ψ, h, φ];

  CloneQureg[ψ, in];
  ApplyCircuit[ψ, u /. v2];
  e2 = CalcExpecPauliString[ψ, h, φ];

  (e1 - e2) / 2
]
```

Let's randomly initialise our parameters to values $v\theta$, and check the corresponding energy.

```
vθ = Table[θ[i] → RandomReal[], {i, nθ}]
initθ = vθ;

{θ[1] → 0.944783, θ[2] → 0.674455, θ[3] → 0.0508029, θ[4] → 0.381874,
 θ[5] → 0.947617, θ[6] → 0.803537, θ[7] → 0.518898, θ[8] → 0.12349,
 θ[9] → 0.00541244, θ[10] → 0.633603, θ[11] → 0.708769, θ[12] → 0.330352,
 θ[13] → 0.84117, θ[14] → 0.300132, θ[15] → 0.315277, θ[16] → 0.38555,
 θ[17] → 0.807733, θ[18] → 0.493529, θ[19] → 0.235609, θ[20] → 0.499855,
 θ[21] → 0.984375, θ[22] → 0.912659, θ[23] → 0.273329, θ[24] → 0.353283,
 θ[25] → 0.399153, θ[26] → 0.549364, θ[27] → 0.983178, θ[28] → 0.889903,
 θ[29] → 0.0928538, θ[30] → 0.975794, θ[31] → 0.807838, θ[32] → 0.16291,
 θ[33] → 0.802887, θ[34] → 0.842839, θ[35] → 0.415662, θ[36] → 0.647286,
 θ[37] → 0.494067, θ[38] → 0.349075, θ[39] → 0.880707, θ[40] → 0.620367,
 θ[41] → 0.453247, θ[42] → 0.000729754, θ[43] → 0.564278, θ[44] → 0.919765,
 θ[45] → 0.607461, θ[46] → 0.248741, θ[47] → 0.784312, θ[48] → 0.838105,
 θ[49] → 0.0971372, θ[50] → 0.889634, θ[51] → 0.152776, θ[52] → 0.461625,
 θ[53] → 0.957948, θ[54] → 0.375586, θ[55] → 0.786615, θ[56] → 0.0641479}

CloneQureg[ψ, in];
ApplyCircuit[ψ, u /. vθ];
CalcExpecPauliString[ψ, h, φ]
0.358099
```

Here are the derivatives $\frac{d}{d\theta[1]} \langle E(\theta) \rangle$ and $\frac{d}{d\theta[2]} \langle E(\theta) \rangle$

```
calcExpecDeriv[h, in, u, vθ, θ[1], ψ, φ]
0.120587
```

```
calcExpecDeriv[h, in, u, vθ, θ[2], ψ, φ]
-0.0609789
```

Obtaining the gradient at parameter values $\mathbf{v}\theta$ therefore requires evaluating $2n\theta$ expectation values.

```
grad = Table[
  calcExpecDeriv[h, in, u, vθ, θ[i], ψ, φ],
  {i, nθ}
]
{0.120587, -0.0609789, -0.0617724, -0.316569, -0.0206436, -0.234134, -0.175736,
 0.25312, 0.330288, -0.156502, -0.129305, -0.103314, -0.189972, 0.0314764,
 0.152959, -0.0179045, 0.279813, -0.352306, -0.0275923, -0.22221, 0.203336,
 -0.021528, 0.309561, -0.060743, 0.139741, -0.154059, 0.234719, -0.266365,
 0.076402, -0.00489835, -0.205105, -0.146627, -0.391673, -0.0854651, 0.103386,
 -0.160257, 0.383993, 0.016442, 0.123013, -0.143761, 0.0446333, -0.383324,
 -0.119879, 0.212571, 0.503745, -0.134012, -0.0179596, -0.16494, 0.32401,
 -0.164754, 0.00892003, 0.0577206, 0.178624, -0.0516252, 0.136484, -0.399077}
```

Gradient descent simply instructs the experimentalist to update the parameters in the opposite direction to this gradient, in order to reduce the cost function. That is:

$$\Delta \vec{\theta} = -\nabla \langle E(\vec{\theta}) \rangle$$

```
Δt = 0.1;
vθ[[All, 2]] -= Δt grad;

CloneQureg[ψ, in];
ApplyCircuit[ψ, u /. vθ];
CalcExpecPauliString[ψ, h, φ]
0.120877
```


Indeed the expectation value has reduced! This process can be repeated, iteratively minimising the cost function.

```
vθ = initθ;
gradVals = Table[
  grad = Table[calcExpecDeriv[h, in, u, vθ, θ[i], ψ, φ], {i, nθ}];
  vθ[[All, 2]] -= Δt grad;

  ApplyCircuit[CloneQureg[ψ, in], u /. vθ];
  CalcExpecPauliString[ψ, h, φ],
  30
];
```



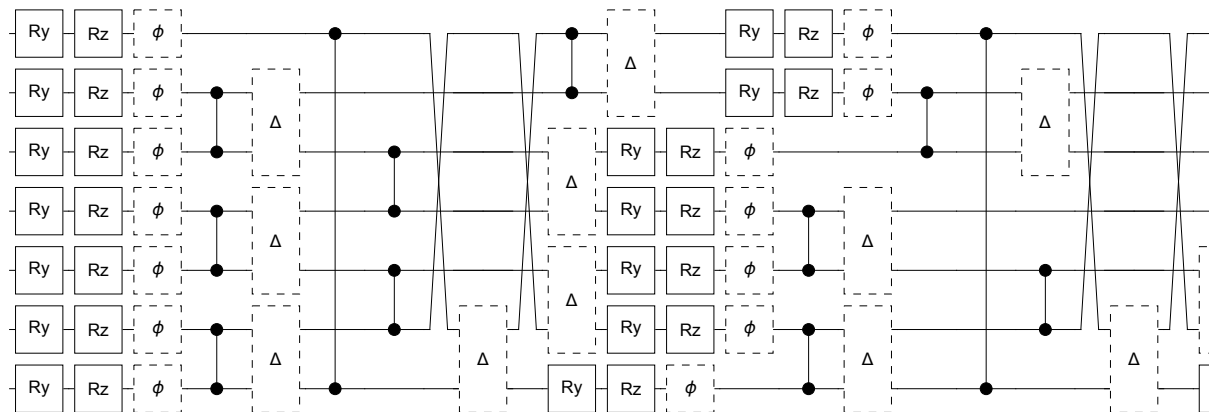
```
ListPlot[gradVals,
  AxesLabel → {"iteration", "energy"},
  PlotRange → {1.1 vMin, .5},
  GridLines → {{}, {{vMin, Directive[Red, Dashed]}}}]
```



We have effectively simulated quantum natural gradient in a noise-free setting. Let's now introduce decoherence into our ansatz circuit, inserting dephasing noise after every R_z , and two-qubit depolarising noise after every control-Z.

```
ch = u /. {
  g : Rz_t[_] => Sequence[g, Deph_t[10^-3]],
  g : Cc_[Z_t_] => Sequence[g, Depol_c,t[10^-2]]};
```

```
DrawCircuit[ch]
```



By now merely changing our states to be density matrices...

```
{ρ, μ, in} = CreateDensityQuregs[nQb, 3];
InitPlusState[in];
```

all our previous calculations can be repeated in the presence of noise.

```
CloneQureg[ $\rho$ , in];
ApplyCircuit[ $\rho$ , ch /. v $\theta$ ];
CalcExpecPauliString[ $\rho$ , h,  $\mu$ ]
-1.9876
```

We can see that the introduced decoherence has damaged the fidelity of the final state.

```
CalcPurity[ $\rho$ ]
0.652987
```


```
CalcFidelity[ $\rho$ ,  $\psi$ ]
0.807163
```

Let's see how the full evolution of gradient descent would differ if noise were present at every iteration.

```
v $\theta$  = init $\theta$ ;
noisyGradVals = Table[
  grad = Table[calcExpecDeriv[h, in, ch, v $\theta$ ,  $\theta$ [i],  $\rho$ ,  $\mu$ ], {i, n $\theta$ ]];
  v $\theta$ [[All, 2]] -=  $\Delta t$  grad;

  ApplyCircuit[CloneQureg[ $\rho$ , in], ch /. v $\theta$ ];
  CalcExpecPauliString[ $\rho$ , h,  $\mu$ ],
  30
];

ListPlot[{gradVals, noisyGradVals},
  AxesLabel → {"iteration", "energy"},
  PlotLegends → {"pure", "noisy"},
  PlotRange → {0, 1.1 vMin},
  GridLines → {{}, {{vMin, Directive[Red, Dashed]}}}]
```



Natural gradient

Superior quantum minimisation techniques exist which converge faster than gradient descent, and more reliably, while requiring only a modest increase in experimental measurements. One

such technique is *natural gradient* which prescribes a change in parameters $\Delta \vec{\theta}$ given by:

$$\text{Re}[\mathbf{G}] (\Delta \vec{\theta}) = -\Delta t \nabla \langle E(\vec{\theta}) \rangle$$

The matrix $\text{Re}[\mathbf{G}]$ is the Fubini-Study metric tensor, equivalent to the real component of the quantum geometric tensor with entries:

$$G_{ij} = \frac{\partial \langle \psi(\vec{\theta}) |}{\partial \theta_i} \frac{\partial | \psi(\vec{\theta}) \rangle}{\partial \theta_j} - \left(\frac{\partial \langle \psi(\vec{\theta}) |}{\partial \theta_i} | \psi(\vec{\theta}) \rangle \right) \left(\langle \psi(\vec{\theta}) | \frac{\partial | \psi(\vec{\theta}) \rangle}{\partial \theta_j} \right)$$

This horrifying looking matrix is thankfully straightforward to evaluate on a quantum computer. And because we have the luxury of *simulating* the algorithm, rather than executing its prescribed circuits on experimental hardware, we can even just directly evaluate the matrix G using:

? CalcMetricTensor

Symbol

`CalcMetricTensor[inQureg, circuit, varVals]` returns the natural gradient metric tensor, capturing the circuit derivatives (produced from initial state `inQureg`) with respect to `varVals`, specified with values `{var -> value, ...}`.

`CalcMetricTensor[inQureg, circuit, varVals, workQuregs]` uses the given persistent workspace quregs (`workQuregs`) in lieu of creating them internally, and should be used for optimum performance. At most four `workQuregs` are needed.

- For state-vectors and pure circuits, this returns the quantum geometric tensor, which relates to the Fubini-Study metric, the classical Fisher information matrix, and the variational imaginary-time Li tensor with Berry connections.
- For density-matrices and noisy channels, this function returns the Hilbert-Schmidt derivative metric, which well approximates the quantum Fisher information matrix, though is a more experimentally relevant minimisation metric (<https://arxiv.org/abs/1912.08660>).
- Variable repetition, multi-parameter gates, variable-dependent element-wise matrices, variable-dependent channels, and operators whose parameters are (numerically evaluable) functions of variables are all permitted.
- All operators must be invertible, trace-preserving and deterministic, else an error is thrown.
- This function runs asymptotically faster than `ApplyCircuitDerivs[]` and requires only a fixed memory overhead.

Let's return to pure-state simulation, and re-randomise the ansatz parameters

```
{ψ, φ, in} = CreateQuregs[nQb, 3];
InitPlusState[in];
```

```

vθ = Table[θ[i] → RandomReal[], {i, nθ}]
{θ[1] → 0.0824625, θ[2] → 0.31652, θ[3] → 0.459314, θ[4] → 0.807564,
 θ[5] → 0.237396, θ[6] → 0.528357, θ[7] → 0.815375, θ[8] → 0.750482,
 θ[9] → 0.830971, θ[10] → 0.698174, θ[11] → 0.740431, θ[12] → 0.0252106,
 θ[13] → 0.594864, θ[14] → 0.154642, θ[15] → 0.16513, θ[16] → 0.240785,
 θ[17] → 0.95795, θ[18] → 0.119302, θ[19] → 0.97886, θ[20] → 0.562465,
 θ[21] → 0.314392, θ[22] → 0.224505, θ[23] → 0.322188, θ[24] → 0.401244,
 θ[25] → 0.288787, θ[26] → 0.137886, θ[27] → 0.805133, θ[28] → 0.338753,
 θ[29] → 0.520798, θ[30] → 0.543042, θ[31] → 0.442366, θ[32] → 0.947235,
 θ[33] → 0.407512, θ[34] → 0.484981, θ[35] → 0.309018, θ[36] → 0.124111,
 θ[37] → 0.620958, θ[38] → 0.945303, θ[39] → 0.816339, θ[40] → 0.455867,
 θ[41] → 0.280547, θ[42] → 0.381771, θ[43] → 0.418341, θ[44] → 0.498481,
 θ[45] → 0.900826, θ[46] → 0.652939, θ[47] → 0.934518, θ[48] → 0.225505,
 θ[49] → 0.133524, θ[50] → 0.742551, θ[51] → 0.290751, θ[52] → 0.969005,
 θ[53] → 0.530504, θ[54] → 0.781994, θ[55] → 0.568455, θ[56] → 0.276424}

```

The Fubini-Study metric tensor is trivially obtained:

```

g = Re @ CalcMetricTensor[in, u, vθ];
g[[50 ;;, 50 ;;]] // Chop // MatrixForm

```

0.249999	0.0006693	0.0103687	-0.0153045	0.00411197	0.00853124	0.0
0.0006693	0.249838	0.000244921	0.00291078	0.00240679	0.0347412	0.0
0.0103687	0.000244921	0.24963	0.00762959	0.00303132	-0.0167092	-0.0
-0.0153045	0.00291078	0.00762959	0.247566	-0.00219198	-0.0134712	0.00
0.00411197	0.00240679	0.00303132	-0.00219198	0.248026	-0.0124202	0.0
0.00853124	0.0347412	-0.0167092	-0.0134712	-0.0124202	0.25	1.293
0.0111326	0.0133983	-0.00565139	0.00640067	0.0142553	1.29301×10^{-7}	0.2

The energy gradient, also required by quantum natural gradient, *can* be obtained in an identical manner to that of quantum gradient descent - via the parameter-shift rule. However, it is more convenient make use of another function below, which uses back-propagation to run $O(\#gates)$ faster!

? CalcExpecPauliStringDerivs**Symbol**

CalcExpecPauliStringDerivs[inQureg, circuit, varVals, pauliString] returns the gradient vector of the pauliString expected values, as produced by the derivatives of the circuit (with respect to varVals, {var -> value}) acting upon the given initial state (inQureg).

CalcExpecPauliStringDerivs[inQureg, circuit, varVals, pauliQureg] accepts a Qureg pre-initialised as a pauli string via SetQuregToPauliString[] to speedup density-matrix simulation.

CalcExpecPauliStringDerivs[inQureg, circuit, varVals, pauliStringOrQureg, workQuregs] uses the given persistent workspaces (workQuregs) in lieu of creating them internally, and should be used for optimum performance. At most four workQuregs are needed.

- Variable repetition, multi-parameter gates, variable-dependent element-wise matrices, variable-dependent channels, and operators whose parameters are (numerically evaluable) functions of variables are all permitted.
- All operators must be invertible, trace-preserving and deterministic, else an error is thrown.
- This function runs asymptotically faster than ApplyCircuitDerivs[] and requires only a fixed memory overhead.

```
grad = CalcExpecPauliStringDerivs[in, u, v0, h]
{0.24538, 0.106511, 0.122928, 0.119399, -0.0220159, 0.128087, -0.0589124,
 0.356504, 0.0371179, -0.20374, 0.350681, 0.0476721, 0.0982792, -0.108923,
 0.248425, 0.10902, -0.219147, -0.0520782, -0.134164, -0.00339067, 0.313054,
 0.36457, -0.0214842, -0.0897153, 0.375175, 0.00480845, 0.35326, -0.0401864,
 -0.143416, -0.0361575, -0.0823219, -0.114158, -0.151913, 0.012574, -0.155273,
 0.391916, -0.240681, -0.00291144, 0.183757, -0.00738578, 0.0687664, -0.00819604,
 0.140992, 0.00546608, 0.193064, -0.346827, 0.0245652, 0.0320816, 0.621942,
 0.339881, -0.148997, -0.0181471, 0.409203, 0.0321777, 0.276417, 0.157329}
```

Now that we have obtained $\text{Re}[G]$ and $\nabla \langle E(\vec{\theta}) \rangle$, computing an iteration of natural gradient requires solving the below linear equation for the change in parameters, $\Delta \vec{\theta}$.

$$\text{Re}[G] (\Delta \vec{\theta}) = -\Delta t \nabla \langle E(\vec{\theta}) \rangle$$

Alas we *cannot* simply invert G to obtain $\Delta \vec{\theta} = -\Delta t \text{Re}[G]^{-1} \nabla \langle E(\vec{\theta}) \rangle$, because it is often singular!

Det [g]

2.919×10^{-72}

Instead, we seek an approximate solution - there are many ways to do this!

```
 $\Delta\theta = \text{LinearSolve}[g, -\Delta t \text{ grad}]$ 
```

```
{-0.465391, 0.664162, -0.528219, -0.492678, -0.0402829, -0.576641, -0.133898,
-0.628493, 0.671918, 0.799776, -0.455777, 0.0848771, 0.476907, -0.0898445,
-0.262464, -1.23912, -0.134213, -0.754336, -0.729321, 0.52531, 0.445998,
0.0802577, -0.114921, 0.602765, 0.445231, -0.212191, -0.285996, 1.66092,
0.435719, -0.355767, 0.614801, 0.780978, 0.353793, -0.547958, 0.157621,
-1.24833, -0.733746, -0.87831, 0.182781, 0.705881, -0.251065, -2.17302,
0.0563383, 0.93814, 0.150695, 0.331929, 0.482414, 0.948237, -0.581249,
1.00339, -0.0569486, 0.209921, -0.438878, -0.367177, 0.119752, 0.582094}
```

```
 $\Delta\theta = -\Delta t \text{ PseudoInverse}[g] \cdot \text{grad}$ 
```

```
{-0.465391, 0.664162, -0.528219, -0.492678, -0.0402829, -0.576641, -0.133898,
-0.628493, 0.671918, 0.799776, -0.455777, 0.0848771, 0.476907, -0.0898445,
-0.262464, -1.23912, -0.134213, -0.754336, -0.729321, 0.52531, 0.445998,
0.0802577, -0.114921, 0.602765, 0.445231, -0.212191, -0.285996, 1.66092,
0.435719, -0.355767, 0.614801, 0.780978, 0.353793, -0.547958, 0.157621,
-1.24833, -0.733746, -0.87831, 0.182781, 0.705881, -0.251065, -2.17302,
0.0563383, 0.93814, 0.150695, 0.331929, 0.482414, 0.948237, -0.581249,
1.00339, -0.0569486, 0.209921, -0.438878, -0.367177, 0.119752, 0.582094}
```

```
 $\Delta\theta = \text{Fit}[\{g, -\Delta t \text{ grad}\}, \text{FitRegularization} \rightarrow \{\text{"Tikhonov"}, 10^{-10}\}]$ 
```

```
{-0.465406, 0.663854, -0.528191, -0.492677, -0.0403096, -0.576662, -0.133886,
-0.628557, 0.671896, 0.799712, -0.45575, 0.0848979, 0.476797, -0.0898166,
-0.262497, -1.23876, -0.134266, -0.75432, -0.729284, 0.525333, 0.445972,
0.0803299, -0.114905, 0.602803, 0.445198, -0.212215, -0.285993, 1.66045,
0.435746, -0.355744, 0.61476, 0.780884, 0.353805, -0.547964, 0.157614,
-1.24801, -0.733727, -0.878255, 0.182781, 0.70579, -0.250976, -2.17244,
0.0564028, 0.938096, 0.150723, 0.332029, 0.482384, 0.948194, -0.581248,
1.0031, -0.0569467, 0.209875, -0.438842, -0.367145, 0.119762, 0.582013}
```

The last method is *Tikhonov regularisation* with parameter $\lambda = 10^{-10}$, which obtains

$\min_{\vec{\theta}} \left| \text{Re}[\mathbf{G}] \left(\Delta \vec{\theta} \right) + \Delta t \nabla \langle E(\vec{\theta}) \rangle \right|^2 + \lambda \left| \Delta \vec{\theta} \right|^2$, additionally constraining the linear equation to minimise the size of the parameter change, enforcing smoothness.

Let's now simulate quantum natural gradient upon the same system we threw at gradient descent.


```

vθ = initθ;
natGradVals = Table[
  grad = CalcExpecPauliStringDerivs[in, u, vθ, h];
  g = Re @ CalcMetricTensor[in, u, vθ];
  Δθ = Fit[{g, -Δt grad}, FitRegularization → {"Tikhonov", 10-6}]];
  vθ[[All, 2]] += Δθ;

  ApplyCircuit[CloneQureg[ψ, in], u /. vθ];
  CalcExpecPauliString[ψ, h, φ],
  30
];

ListPlot[{gradVals, natGradVals},
  AxesLabel → {"iteration", "energy"},
  PlotLegends → {"gradient descent", "natural gradient"},
  PlotRange → {0, 1.1 vMin},
  GridLines → {{}, {{vMin, Directive[Red, Dashed]}}}]]

```



These functions are also compatible with density matrices and parameterised channels!

```

{ρ, μ, in} = CreateDensityQuregs[nQb, 3];
InitPlusState[in];


vθ = initθ;
noisyNatGradVals = Table[
  grad = CalcExpecPauliStringDerivs[in, ch, vθ, h];
  g = Re @ CalcMetricTensor[in, ch, vθ];
  Δθ = Fit[{g, -Δt grad}, FitRegularization → {"Tikhonov", 10-6}]];
  vθ[[All, 2]] += Δθ;

  ApplyCircuit[CloneQureg[ρ, in], ch /. vθ];
  CalcExpecPauliString[ρ, h, μ],
  30
];

```

Let's compare all our simulations.

```
ListPlot[{gradVals, noisyGradVals, natGradVals, noisyNatGradVals},
  AxesLabel → {"iteration", "energy"},
  PlotLegends → {
    "pure gradient descent",
    "noisy gradient descent",
    "pure natural gradient",
    "noisy natural gradient"},
  PlotRange → {0, 1.1 vMin},
  GridLines → {{}, {{vMin, Directive[Red, Dashed]}}}]
```



Note this is not an interpretable performance comparison of these methods; we have not individually optimised the timesteps Δt , nor compared their resource costs. We have made no effort to sensibly initialise our parameters, avoid barren plateaus, nor reason about convergence.

We finally mention that watching the parameters smoothly evolve during minimisation can be quite a show!

```
in = InitPlusState @ CreateQureg[nQb];
vθ = Table[θ[i] → RandomReal[], {i, nθ}];
θt = Transpose @ Table[
  grad = CalcExpecPauliStringDerivs[in, u, vθ, h];
  g = Re @ CalcMetricTensor[in, u, vθ];
  Δθ = Fit[{g, -Δt grad}, FitRegularization → {"Tikhonov", 10-6}]];
  vθ[[All, 2]] += Δθ,
  30
];
```



```
ListLinePlot[ $\theta$ t, PlotRange → {{1, All}, All}, PlotStyle → "DeepSeaColors"]
```

